

CS302
Operating System
Lab 4

Thread and Schedule


ZENG Xinxun, ZHANG Shiqi

Outline

- Thread
 - Why multithreading?
 - Thread libraries
 - Practice
 - Threading issue
- Schedule
 - Why scheduling?
 - Three level scheduling
- Assignment: Scheduler

Why multithreading?


- Responsiveness
- Scalability
- Resource sharing
- Economy



Picture from: <https://computing.llnl.gov/tutorials/pthreads/>

Threading vs Processing

Platform	fork()			pthread_create()		
	real	user	sys	real	user	sys
Intel 2.6 GHz Xeon E5-2670 (16 cores/node)	8.1	0.1	2.9	0.9	0.2	0.3
Intel 2.8 GHz Xeon 5660 (12 cores/node)	4.4	0.4	4.3	0.7	0.2	0.5
AMD 2.3 GHz Opteron (16 cores/node)	12.5	1.0	12.5	1.2	0.2	1.3
AMD 2.4 GHz Opteron (8 cores/node)	17.6	2.2	15.7	1.4	0.3	1.3
IBM 4.0 GHz POWER6 (8 cpus/node)	9.5	0.6	8.8	1.6	0.1	0.4
IBM 1.9 GHz POWER5 p5-575 (8 cpus/node)	64.2	30.7	27.6	1.7	0.6	1.1
IBM 1.5 GHz POWER4 (8 cpus/node)	104.5	48.6	47.2	2.1	1.0	1.5
INTEL 2.4 GHz Xeon (2 cpus/node)	54.9	1.5	20.8	1.6	0.7	0.9
INTEL 1.4 GHz Itanium2 (4 cpus/node)	54.5	1.1	22.2	2.0	1.2	0.6

 [fork_vs_thread.txt](#)

Picture from: <https://computing.llnl.gov/tutorials/pthreads/>

Time utility of 50,000 process/thread creation

Thread libraries

- Three main libraries:
 - POSIX Pthread
 - Windows thread
 - Java thread

POSIX Pthread

- Compile flag: -pthread e.g. gcc thread.c -o thread -pthread
- API
 - pthread_create
 - pthread_exit
 - pthread_cancel
 - pthread_join
 - pthread_detach

```
6  struct msg{
7      int id;
8      char* word;
9  };
10
11 void PrintHello(struct msg* arg)
12 {
13     printf("%s%d!\n", arg->word, arg->id);
14     pthread_exit(NULL);
15 }
16
17 int main(int argc, char *argv[])
18 {
19     pthread_t threads[NUM_THREADS];
20     struct msg arg[NUM_THREADS];
21     char* word = "Hello World! It's me, thread #";
22     for(int t=0;t<NUM_THREADS;t++){
23         printf("In main: creating thread %d\n", t);
24         arg[t] = (struct msg){t, word};
25         pthread_create(&threads[t], NULL, &PrintHello, &arg[t]);
26     }
27     for(int t=0;t<NUM_THREADS;t++)
28         pthread_join(threads[t], NULL);
29     printf("\nall threads finish.\n");
30     /* Last thing that main() should do */
31     pthread_exit(NULL);
32 }
```

Implicit Threading

- Thread pool
 - A management pattern of threads.
- OpenMP
 - Threading automatically by compiler.
 - Try: `gcc openmp1.c -o openmp1 -fopenmp`
 - Try: `gcc openmp2.c -o openmp2 -fopenmp`

```
5  int main (int argc, char *argv[])
6  {
7      printf("Number of threads = %d\n", omp_get_num_threads());
8      #pragma omp parallel
9      {
10         printf("I am thread %d\n", omp_get_thread_num());
11     }
12 }
```

```
5  int main (int argc, char *argv[])
6  {
7      int a[80],b[80],c[80];
8      for(int i=0;i<80;++i)
9      {
10         a[i] = i;
11         b[i] = 2*i;
12     }
13
14     #pragma omp parallel for
15     for (int i = 0; i < 80; ++i)
16     {
17         c[i] = a[i] + b[i];
18         printf("i=%d, c[i]=%d\n", i, c[i]);
19     }
20 }
```

Practice

- Write a multithreaded program, user input a series of number, calculates three statistical values in different thread, print the result in given order.
- Input: 90 81 78 95 79 72 85
- Output: average: 82
 minimum: 72
 maximum: 95

Threading issue

- `fork()` and `exec()`
- Signal handling
- Thread cancellation
- Thread-local storage
- Scheduler activations

In Operating System Concept – 4.6

Overview


- Program, Process, Thread, Job
- OS Scheduler:
 - Long Term, Medium Term and Short Term
- Assignment: Scheduler

Program, Process, Thread, Job

- **Program:** a group of instructions to carry out a specified task
 - stored on disk in some file.
- **Process:** an execution of certain program with its own address space
 - resources: CPU, memory address, disk, I/O etc.
- **Thread:** the unit of execution in CPU
- **Job:** a series of program submitted to operating system for some goals.

OS Scheduling

- **A job, after being submitted to operating system, will experience three kind of scheduling.**
 - Long Term Scheduler
 - Medium Term Scheduler
 - Short Term Scheduler

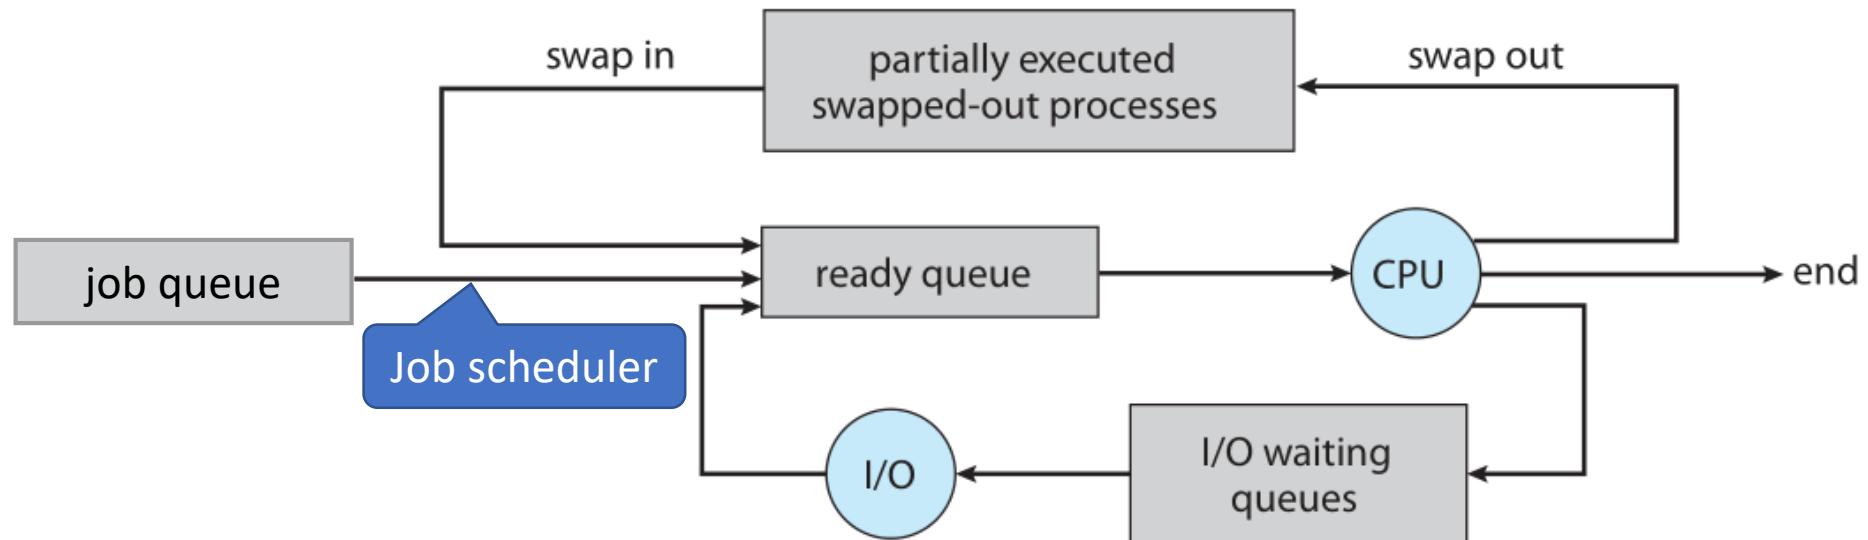


Background Terminology

- **Degree of multiprogramming**
 - the number of processes in memory
- **I/O-bound process**
 - spends more time doing I/O than doing computations
- **CPU-bound process**
 - doing I/O infrequently but using more time on computations

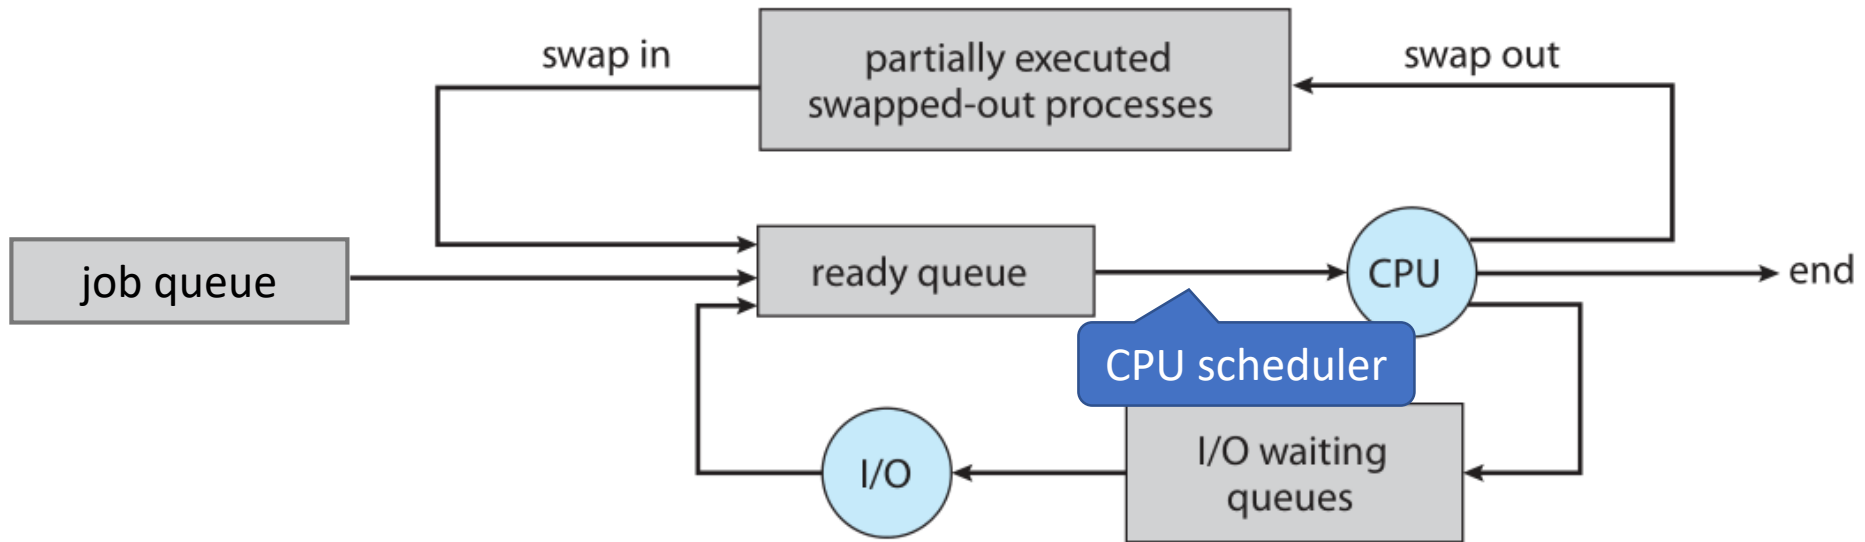
OS Scheduler

- Long Term Scheduler (**Job scheduler**)
 - Select processes from job queue(disk)
 - Load processes to ready queue(main memory)
 - Select a good **process mix** of *IO-bound* and *CPU-bound* processes
 - Control the **degree of multiprogramming**



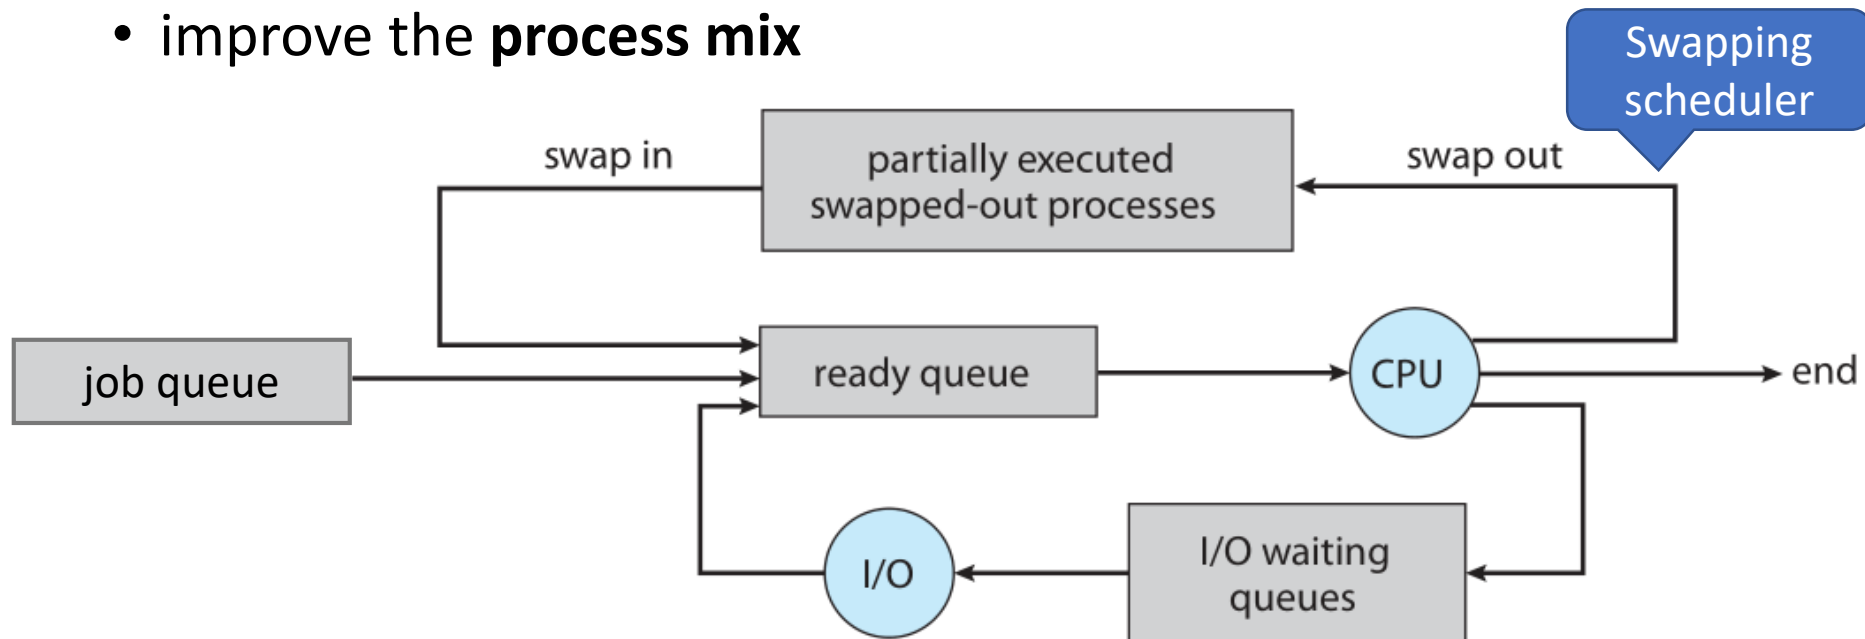
OS Scheduler

- Short Term Scheduler (**CPU scheduler**)
 - Select one process from ready queue
 - Execute it on CPU
 - Scheduling Algorithms: SJF, Round-Robin, Priority

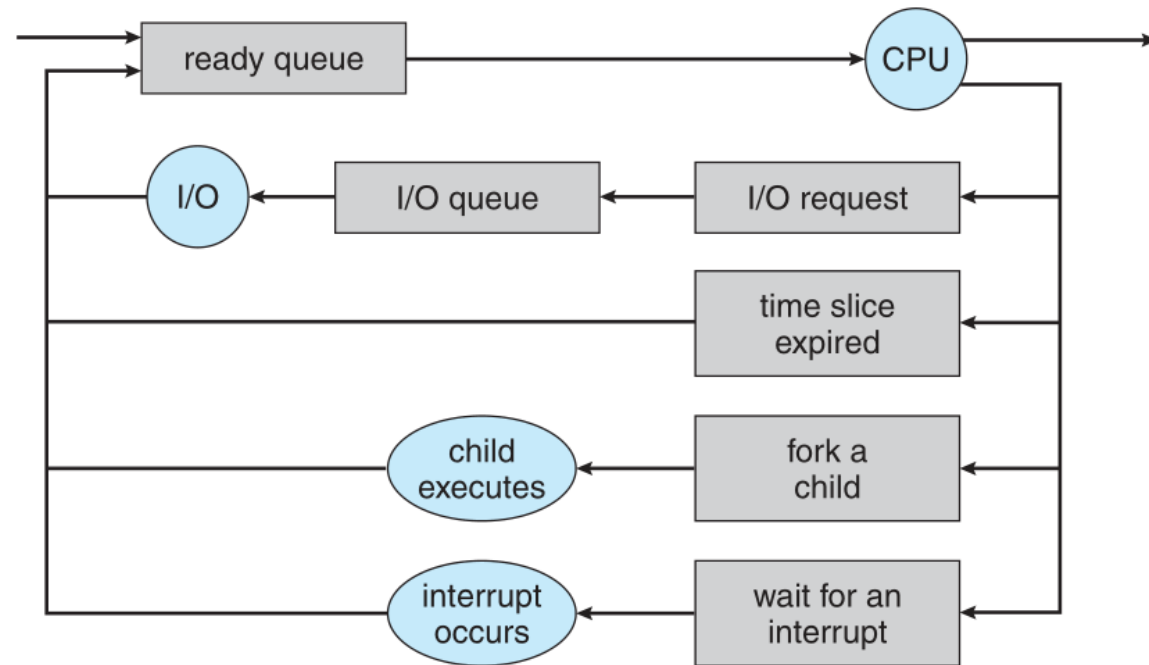


OS Scheduler

- Medium Term Scheduler(**Swapping scheduler**)
 - **swap** a process(IO-bound/CPU-bound) **out** of memory
 - **swap** the process back **into** memory
 - reduce the **degree of multiprogramming**
 - improve the **process mix**



Where a process go after execution?



- normal terminate
- issue an I/O request and then be placed in an I/O queue
- create a new child process and wait for the child's termination
- remove from CPU due to an interrupt and back in the ready queue

OS Scheduler Comparison

Long-Term Scheduler	Short-Term Scheduler	Medium-Term Scheduler
It is a job scheduler	It is a CPU scheduler	It is a process swapping scheduler
Speed is lesser than short term scheduler	Speed is fastest among other two	Speed is in between both short and long term scheduler
It controls the degree of multiprogramming	It provides lesser control over degree of multiprogramming	It reduces the degree of multiprogramming
It selects processes from pool(disk) and loads them into memory for execution	It selects those processes which are ready to execute	It can re-introduce the process into memory and execution can be continued

Assignment

- You are asked to finish a simple CPU scheduler.
- One can add, remove, query jobs, the scheduler can schedule the job according the requirements.

Strategy of CPU scheduler

- Nonpreemptive
- Preemptive
 - Priority
 - Shortest job first
 - Round-Robin

Background

- Strategy: Priority based round-robin, time slice: 100ms
- Job Status: READY, RUNNING, DONE
- Priority Level: 0, 1, 2, 3 (from low to high)
- Initial priority vs Current priority

Assignment

- File organization:
 - **sample.c**: a sample job which runs more than 100ms
 - **scheduler.c**: the main part of this scheduling system
 - **enq.c**: operations about job enqueue
 - **deq.c**: operations about job dequeue
 - **stat.c**: operations about printing status
 - **job.h**: relevant data structures for job

Scheduler.c

- Function:
 - Create a process for a new job. Set its state to READY and put it in the waiting queue.
 - A job will dequeue when receive dequeue request, and the relevant data structure will be cleared.
 - If the job is currently running, it will first stop running and then dequeue.
 - If status request, it outputs information about the currently running job and all jobs in the waiting queue

Enq.c

- Send a enqueue request to the scheduler and submit the job for running.
- Scheduler
 - assigns a unique jid to each job;
 - create a process for it and set its status to READY;
 - put this process into the waiting queue.

```
enq [-p num] e_file args
```

```
-p num:  optional, assign an initial priority  
e_file: the name of the executable file  
args:   the args of the executable file
```


Deq.c

- Send a dequeue request to the scheduler

deq jid

jid: jid that need dequeue

stats.c

- Print out the following information of current running jobs and jobs in waiting queue on standard output:
 - pid
 - user name
 - execution time
 - waiting time
 - create time
 - job status

Job.h

- Define the data struct of all things.
- Define the path of FIFO(use for communicate between all components). You need to create a fifo file in this path to let it work.

Task

- **Modify** `do_stat()` in `scheduler.c` such that it can also output ***jobname, current priority*** and ***default priority***
- **Implement schedule strategy:** job with highest priority runs first, if with same priority, the job waiting longest runs first
 - Priority add 1 after waiting for 100ms
 - Priority reset to initial priority after running
- **Run** this scheduler system
 - do enqueue, dequeue and status operations
 - answer the questions in report