


Lecture 7: Deadlock

Bo Tang @ 2021, Spring

Deadlock

I won't give you the paper unless you come out

I won't come out unless you give me the paper




waiting

holding

waiting




Starvation vs Deadlock



◆ Starvation vs. Deadlock



- ◆ Starvation: thread waits indefinitely
 - ◆ Low-priority thread waiting for resources constantly in use by high-priority threads
- ◆ Deadlock: circular waiting for resources
 - ◆ Thread A owns Res 1 and is waiting for Res 2
 - ◆ Thread B owns Res 2 and is waiting for Res 1



- ◆ Deadlock \Rightarrow Starvation but not vice versa
 - ◆ Starvation can end (but does not have to)
 - ◆ Deadlock cannot end without external intervention


Conditions for Deadlock

- ◆ Deadlock not always deterministic

| <u>process A</u> | <u>process B</u> |
|------------------|------------------|
| sem_wait(x); | sem_wait(y); |
| sem_wait(y); | sem_wait(x); |
| sem_post(y); | sem_wait(x); |
| sem_post(x); | sem_wait(y); |

- ◆ Deadlock will not always happen with this code
 - ◆ Have to have exactly the right timing
 - ◆ So you release a piece of software, and you tested it, and there it is, controlling a nuclear power plant...
- ◆ Deadlocks occur with multiple resources
 - ◆ Means you cannot decompose the problem
 - ◆ Cannot solve deadlock for each resource independently
 - ◆ System with 2 disk drives and two threads
 - ◆ Each process needs 2 disk drives to function
 - ◆ Each process gets one disk and waits for another one


Bridge Crossing Example



- ◆ Each segment of road can be viewed as a resource
 - ◆ Car must own the segment under them
 - ◆ Must acquire segment that they are moving into
- ◆ For bridge: must acquire both halves
 - ◆ Traffic only in one direction at a time
 - ◆ Problem occurs when two cars in opposite directions on bridge: each acquires one segment and needs next
- ◆ If a deadlock occurs, it can be resolved if one car backs up (preempt resources and rollback)
 - ◆ Several cars may have to be backed up
- ◆ Starvation is possible
 - ◆ East-going traffic really fast \Rightarrow no one goes west


Train Example

- ◇ Circular dependency (Deadlock!)
- ◇ Each train wants to turn right
- ◇ Blocked by other trains
- ◇ Similar problem to multiprocessor networks



Train Example

- ◆ Fix? Imagine grid extends in all four directions
 - ◆ **Force ordering of channels** (tracks)
 - ◆ Protocol: Always go east-west first, then north-south
 - ◆ Called “dimension ordering” (X then Y)



Four requirements for Deadlock

◆ Mutual exclusion

- ◆ Only one thread at a time can use a resource.

◆ Hold and wait

- ◆ Thread holding at least one resource is waiting to acquire additional resources held by other threads

◆ No preemption

- ◆ Resources are released only voluntarily by the thread holding the resource, after thread is finished with it


◆ Circular wait

- ◆ There exists a set $\{P_1, \dots, P_n\}$ of waiting processes
 - ◆ P_1 is waiting for a resource that is held by P_2
 - ◆ P_2 is waiting for a resource that is held by P_3
 - ◆ ...
 - ◆ P_n is waiting for a resource that is held by P_1

Resource-Allocation Graph

◆ System Model

- ◆ A set of Threads P_1, P_2, \dots, P_n
- ◆ Resource types R_1, R_2, \dots, R_m
CPU cycles, memory space, I/O devices
- ◆ Each resource type R_i has W_i instances
- ◆ Each thread utilizes a resource as follows:
 - ◆ Request () / Use () / Release ()




◆ Resource-Allocation Graph:

- ◆ V is partitioned into two types:
 - ◆ $P = \{P_1, P_2, \dots, P_n\}$, the set processes in the system.
 - ◆ $R = \{R_1, R_2, \dots, R_m\}$, the set of resource types in system
- ◆ request edge – directed edge $P_1 \rightarrow R_j$
- ◆ assignment edge – directed edge $R_j \rightarrow P_i$


Resource Allocation Graph Examples

◆ Recall:


- ◆ request edge – directed edge $P_1 \rightarrow R_j$
- ◆ assignment edge – directed edge $R_j \rightarrow P_i$



Simple Resource
Allocation Graph




Allocation Graph
With Deadlock



Allocation Graph with
Cycle, but **No Deadlock**

Methods for Handling Deadlocks




- ◆ Allow system to enter deadlock and then recover
 - ◆ Requires deadlock detection algorithm
 - ◆ Some technique for forcibly preempting resources and/or terminating tasks
- ◆ Ensure that system will *never* enter a deadlock
 - ◆ Need to monitor all lock acquisitions
 - ◆ Selectively deny those that *might* lead to deadlock
- ◆ Ignore the problem and pretend that deadlocks never occur in the system
 - ◆ Used by most operating systems, including UNIX

Deadlock Detection Algorithm

- ◆ Only one of each type of resource \Rightarrow look for loops
- ◆ More general deadlock detection algorithm
 - ◆ Let $[X]$ represent an m-ary vector of non-negative integers (quantities of resources of each type):
 - ◆ $[FreeResources]$: current free resources each type
 - ◆ $[Request_x]$: current requests from process X
 - ◆ $[Alloc_x]$: current resources held by process X
 - ◆ See if tasks can eventually terminate on their own

```
[Avail] = [FreeResources]
Add all nodes to UNFINISHED
do {
    done = true
    Foreach node in UNFINISHED {
        if ([Requestnode] <= [Avail]) {
            remove node from UNFINISHED
            [Avail] = [Avail] + [Allocnode]
            done = false
        }
    }
} until(done)
```

- ◆ Nodes left in UNFINISHED \Rightarrow deadlocked



What to do when detect deadlock?

- ◆ Terminate thread, force it to give up resources
 - ◆ In Bridge example, Godzilla picks up a car, throw it into the river. Deadlock solved!
 - ◆ Shoot a dining philosopher
 - ◆ But, not always possible – killing a process holding a mutex leaves world inconsistent
- ◆ Preempt resources without killing off process
 - ◆ Take away resources from process temporarily
 - ◆ Does not always fit with semantics of computation
- ◆ Roll back actions of deadlocked threads
 - ◆ For bridge example, make one car roll backwards (may require others behind him)
 - ◆ Common technique in databases (transactions)
 - ◆ Of course, if you restart in exactly the same way, may reenter deadlock once again
- ◆ Many operating systems use other options

Techniques for Preventing Deadlock

◆ Infinite resources

- ◆ Include enough resources so that no one ever runs out of resources. Examples:
 - ◆ Bay bridge with 12,000 lanes. Never wait!
 - ◆ Infinite disk space (not realistic yet?)

◆ No sharing of resources (totally independent threads)

- ◆ Not very realistic

◆ Do not allow waiting

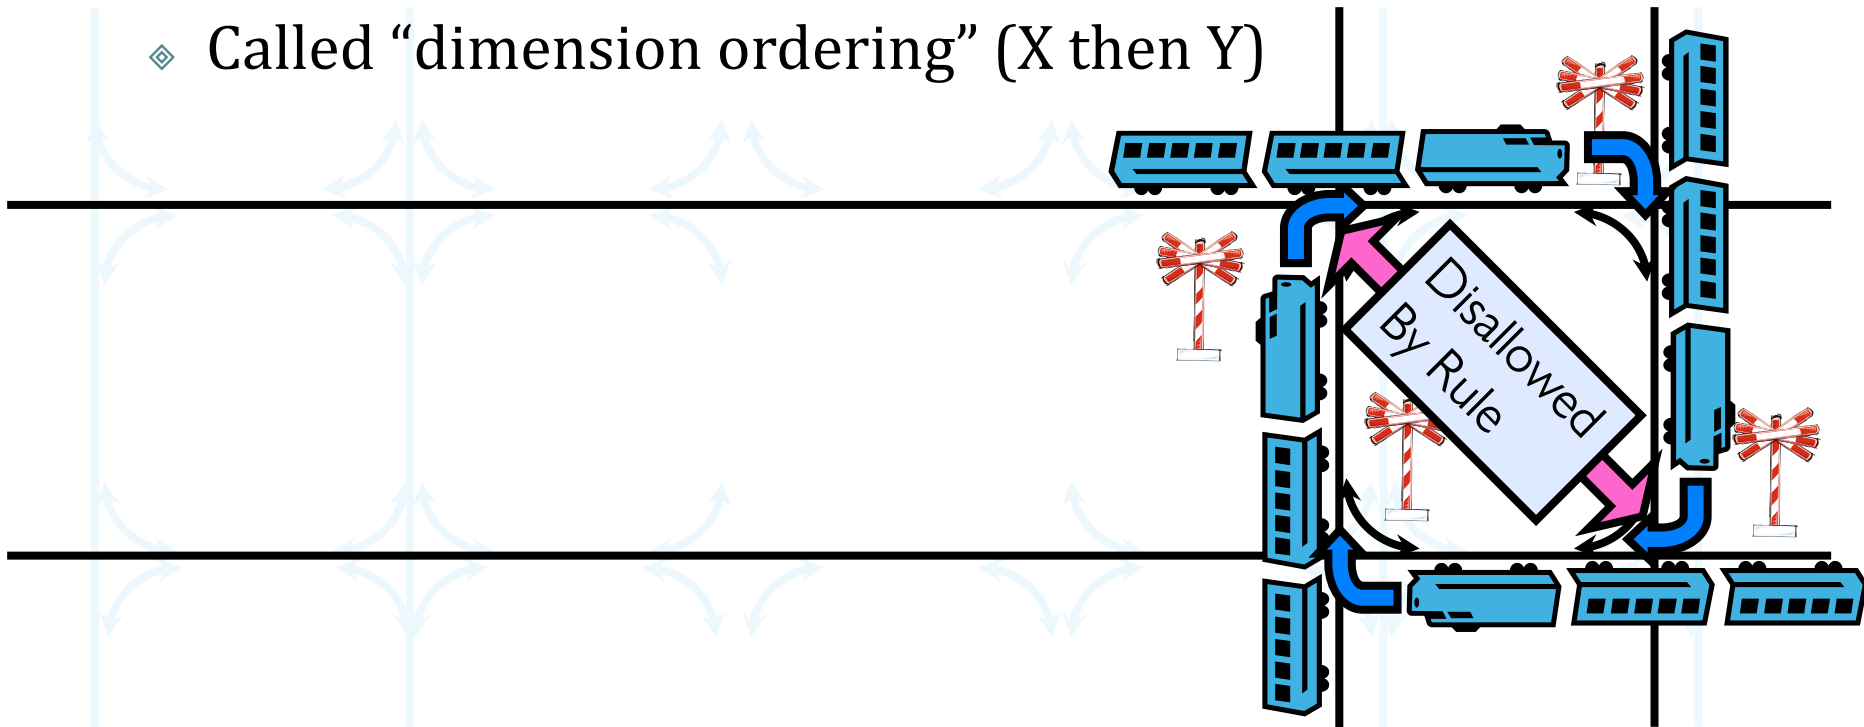
- ◆ Technique used in Ethernet/some multiprocessor nets
 - ◆ Everyone speaks at once. On collision, back off and retry
- ◆ Inefficient, since have to keep retrying
 - ◆ Consider: driving to SUSTech; when hit traffic jam, suddenly you are transported back home and told to retry!

Techniques for Preventing Deadlock

- ◆ Make all processes request everything they will need at the beginning.
 - ◆ Problem: Predicting future is hard, tend to over-estimate resources. Example:
 - ◆ If need 2 chopsticks, request both at same time
 - ◆ Do not leave home until we know no one is using any intersection between home and SUSTech; only one car on the Bay Bridge at a time
- ◆ Force all processes to request resources in a particular order preventing any cyclic use of resources
 - ◆ Thus, preventing deadlock
 - ◆ Example (x.P, y.P, z.P,...)
 - ◆ Make tasks request disk, then memory, then...
 - ◆ Keep from deadlock on freeways around SF by requiring everyone to go clockwise

Review: Train Example

- ◆ Circular dependency (Deadlock!)
 - ◆ Each train wants to turn right, blocked by other trains
 - ◆ Similar problem to multiprocessor networks
- ◆ Fix? Imagine grid extends in all four directions
 - ◆ Force ordering of channels (tracks)
 - ◆ Protocol: Always go east-west first, then north-south
 - ◆ Called “dimension ordering” (X then Y)



Banker's Algorithm for Preventing Deadlock

◆ Toward right idea:

- ◆ State maximum (max) resource needs in advance
- ◆ Allow particular process to proceed if:
 - ◆ $(\text{available resources} - \text{\#requested}) \geq \text{max}$
Remaining that might be needed by any process



◆ Banker's algorithm:

- ◆ Allocate resources dynamically
 - ◆ Evaluate each request and grant if some ordering of processes is still deadlock free afterward
 - ◆ Technique: pretend each request is granted, then run deadlock detection algorithm, substituting $([\text{Maxnode}] - [\text{Allocnode}] \leq [\text{Avail}])$ for $([\text{Requestnode}] \leq [\text{Avail}])$
Grant request if result is deadlock free (conservative!)

(a) [8 pts] Consider the following snapshot of a system

| | Allocation | | | | Max | | | | Available | | | |
|----|------------|---|---|---|-----|---|---|---|-----------|---|---|---|
| | A | B | C | D | A | B | C | D | A | B | C | D |
| P0 | 0 | 0 | 1 | 2 | 0 | 0 | 1 | 2 | 1 | 5 | 2 | 0 |
| P1 | 1 | 0 | 0 | 0 | 1 | 7 | 5 | 0 | | | | |
| P2 | 1 | 3 | 5 | 4 | 2 | 3 | 5 | 6 | | | | |
| P3 | 0 | 0 | 1 | 4 | 0 | 6 | 5 | 6 | | | | |

Answer the following questions using the banker's algorithm:

1. What is the content of the matrix **Need**?
2. Is the system in a safe state? If yes, please give a safe sequence.
3. If a request from process P1 arrives for (0,4,2,0), can the request be granted immediately?
Please explain your answer.

Banker's Algorithm for Preventing Deadlock

```
◆ [Avail] = [FreeResources]
  Add all nodes to UNFINISHED
  do {
    done = true
    Foreach node in UNFINISHED {
      if ([Requestnode] <= [Avail]) {
        remove node from UNFINISHED
        [Avail] = [Avail] + [Allocnode]
        done = false
      }
    }
  } until(done)
```

◆ Allocate resources dynamically

- ◆ Evaluate each request and grant if some ordering of threads is still deadlock free afterward
- ◆ **Technique: pretend each request is granted, then run deadlock detection algorithm, substituting $([Max_{node}] - [Alloc_{node}] \leq [Avail])$ for $([Request_{node}] \leq [Avail])$**
Grant request if result is deadlock free (conservative!)



Banker's Algorithm for Preventing Deadlock

```
[Avail] = [FreeResources]
Add all nodes to UNFINISHED
do {
    done = true
    Foreach node in UNFINISHED {
        if ([Requestnode] <= [Avail]) {
            remove node from UNFINISHED
            [Avail] = [Avail] + [Allocnode]
            done = false
        }
    }
} until(done)
```

ordering of threads is still deadlock free afterward

```
[Avail] = [FreeResources]
Add all nodes to UNFINISHED
do {
    done = true
    Foreach node in UNFINISHED {
        if ([Maxnode] - [Allocnode] <= [Avail]) {
            remove node from UNFINISHED
            [Avail] = [Avail] + [Allocnode]
            done = false
        }
    }
} until(done)
```



Banker's Algorithm for Preventing Deadlock

- ◆ Toward right idea:


- ◆ State maximum resource needs in advance
- ◆ Allow particular process to proceed if:
(available resources - #requested) \geq max
remaining that might be needed by any process

- ◆ Banker's algorithm (less conservative):

- ◆ Allocate resources dynamically
 - ◆ Evaluate each request and grant if some ordering of threads is still deadlock free afterward
 - ◆ **Technique: pretend each request is granted, then run deadlock detection algorithm, substituting $([Max_{node}] - [Alloc_{node}] \leq [Avail])$ for $([Request_{node}] \leq [Avail])$**
Grant request if result is deadlock free (conservative!)
 - ◆ Keeps system in a "SAFE" state, i.e. there exists a sequence $\{P_1, P_2, \dots, P_n\}$ with P_1 requesting all remaining resources, finishing, then P_2 requesting all remaining resources, etc..
- ◆ Algorithm allows the sum of maximum resource needs of all current threads to be greater than total resources



Banker's Algorithm Example



- ◆ Banker's algorithm with dining philosophers
 - ◆ “Safe” (will not cause deadlock) if when try to grab chopstick either:
 - ◆ Not last chopstick
 - ◆ Is last chopstick but someone will have two afterwards
 - ◆ What if k-handed philosophers? Do not allow if:
 - ◆ It is the last one, no one would have k
 - ◆ It is 2nd to last, and no one would have k-1
 - ◆ It is 3rd to last, and no one would have k-2
 - ◆ ...



Thank You!

14. In a single-processor multi-process system, if there are several ready processes, then the wrong statement about process scheduling is ().
- A. process scheduling can execute when a process terminates
 - B. process scheduling can execute when a process creates
 - C. process scheduling can execute when a process is in critical section
 - D. process scheduling can execute when system call finished and system returns to user space